


CS 160 Compilers


Lecture 14.2: Code Generation & llvmLite

Yu Feng
Spring 2023

A typical flow of a compiler



Chomsky hierarchy

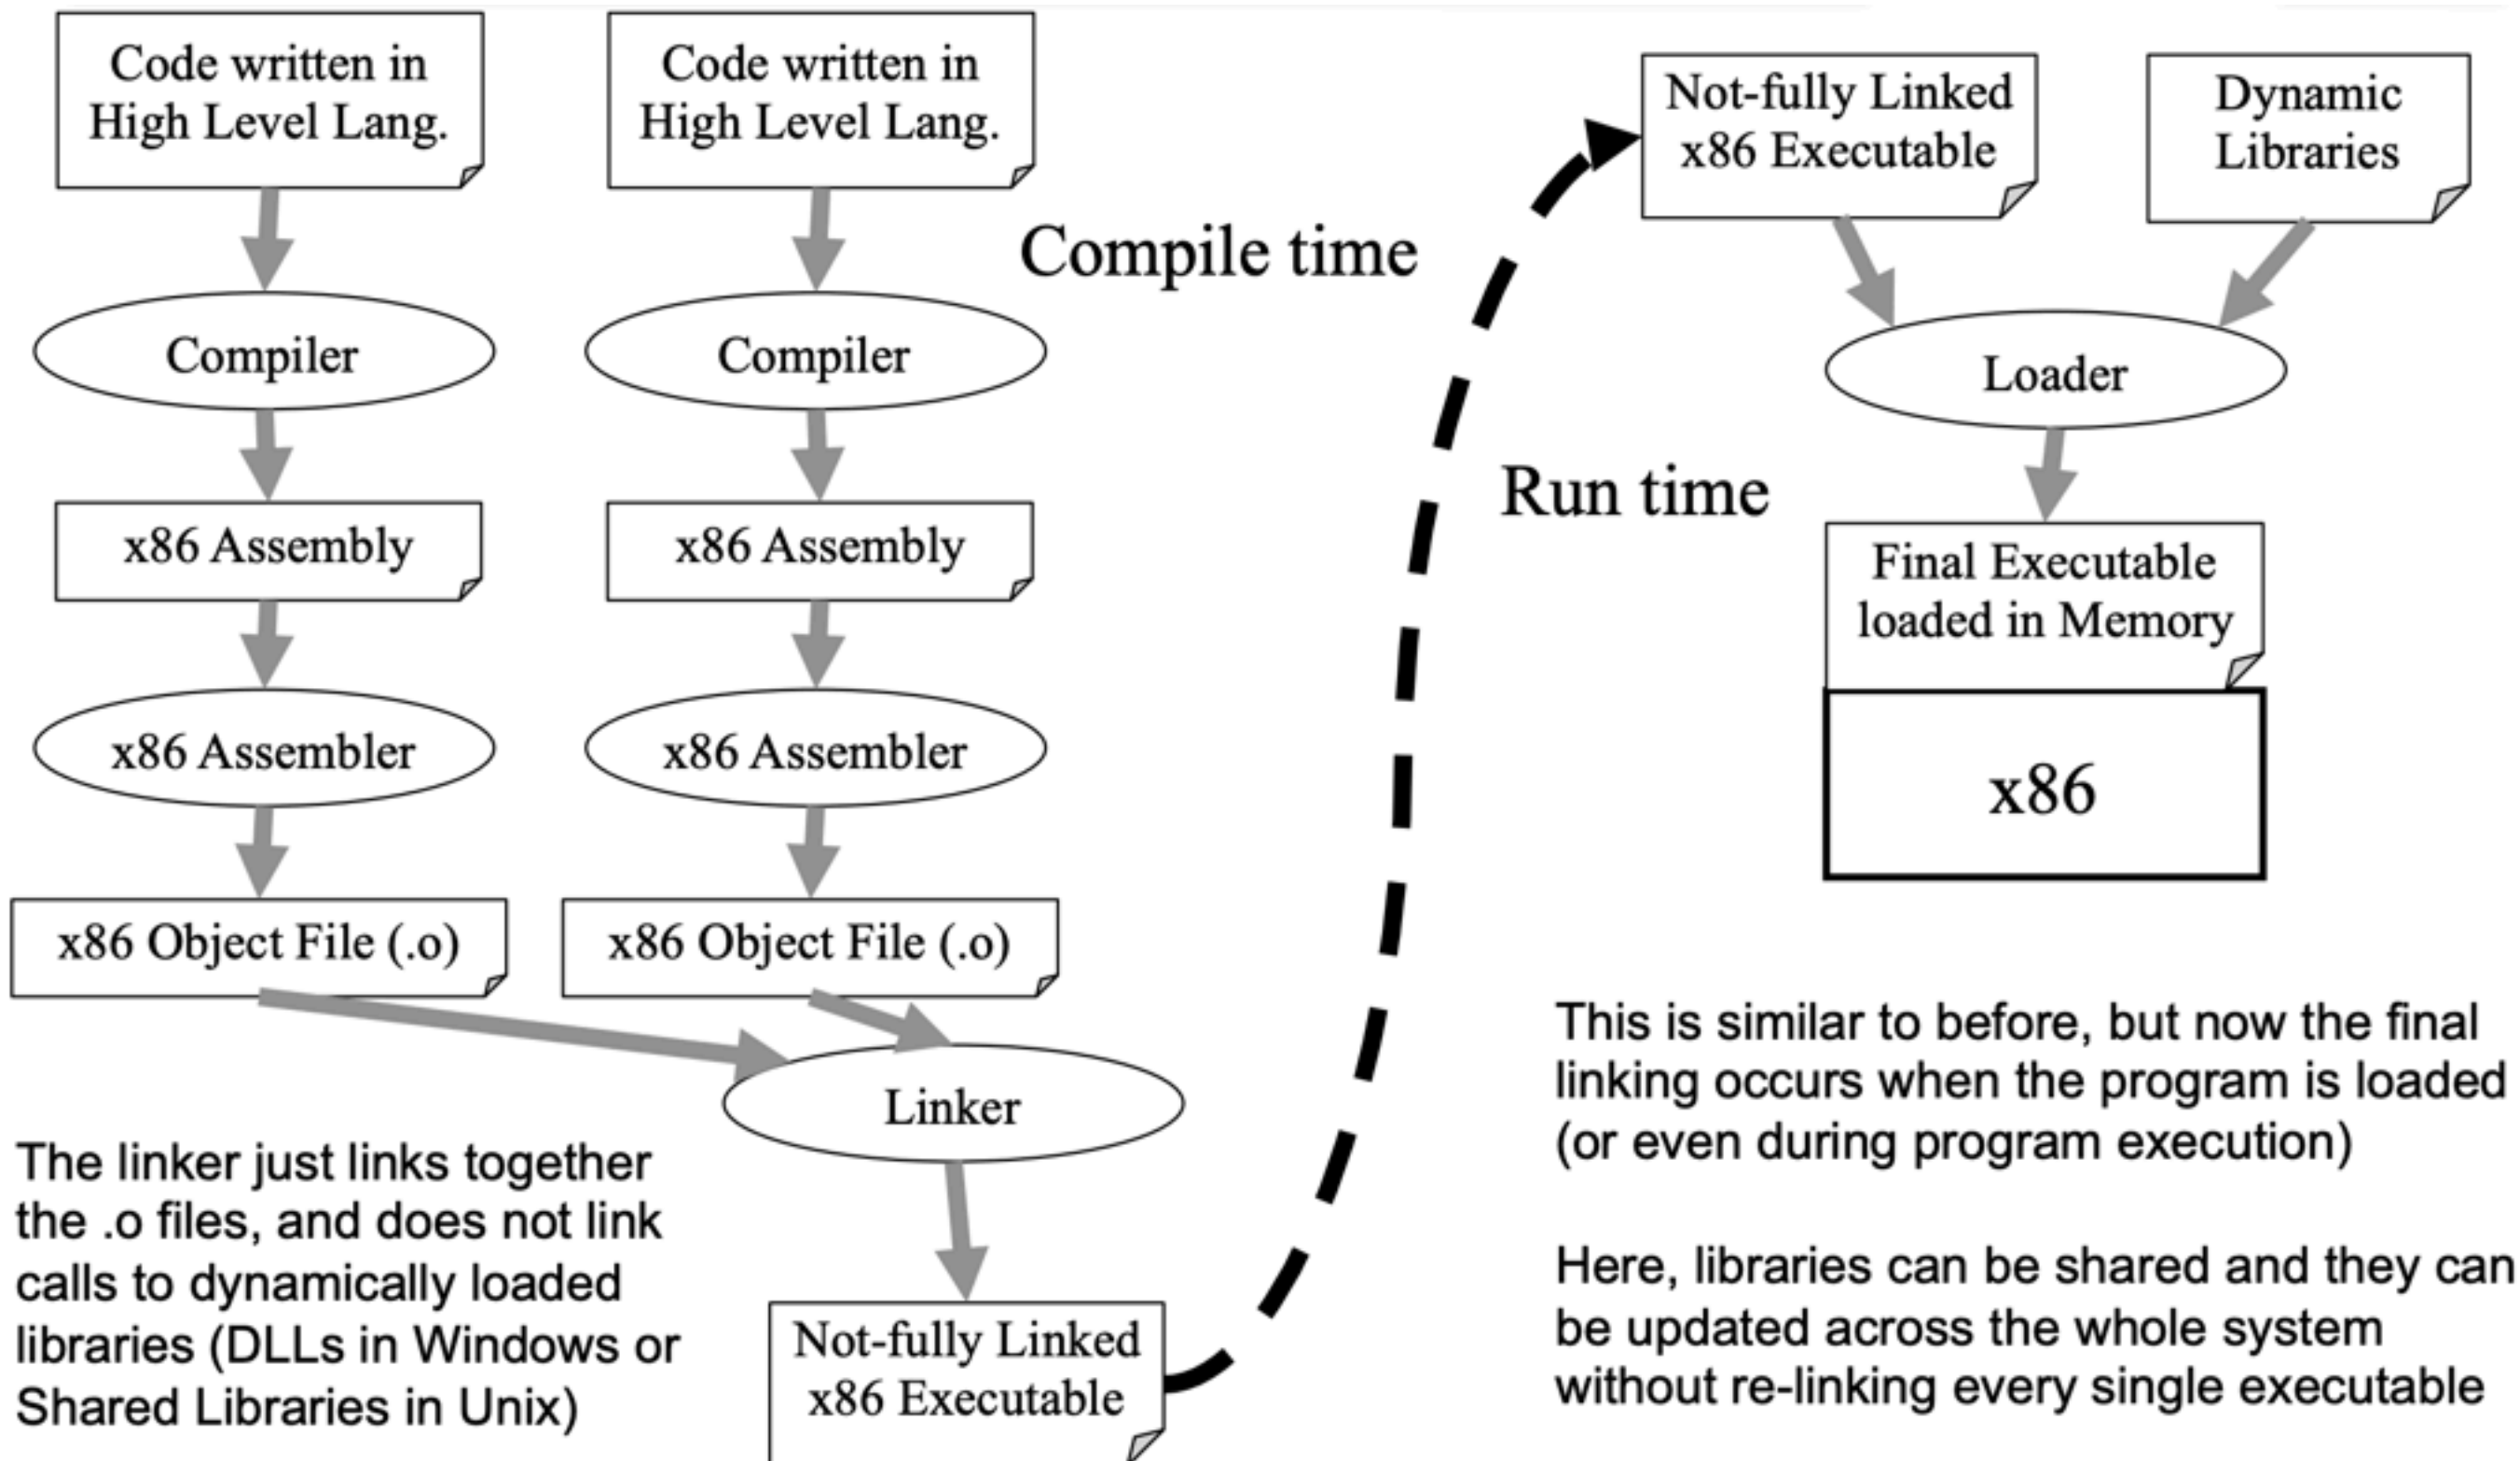


<https://en.wikipedia.org/wiki/File:Chomsky-hierarchy.svg>

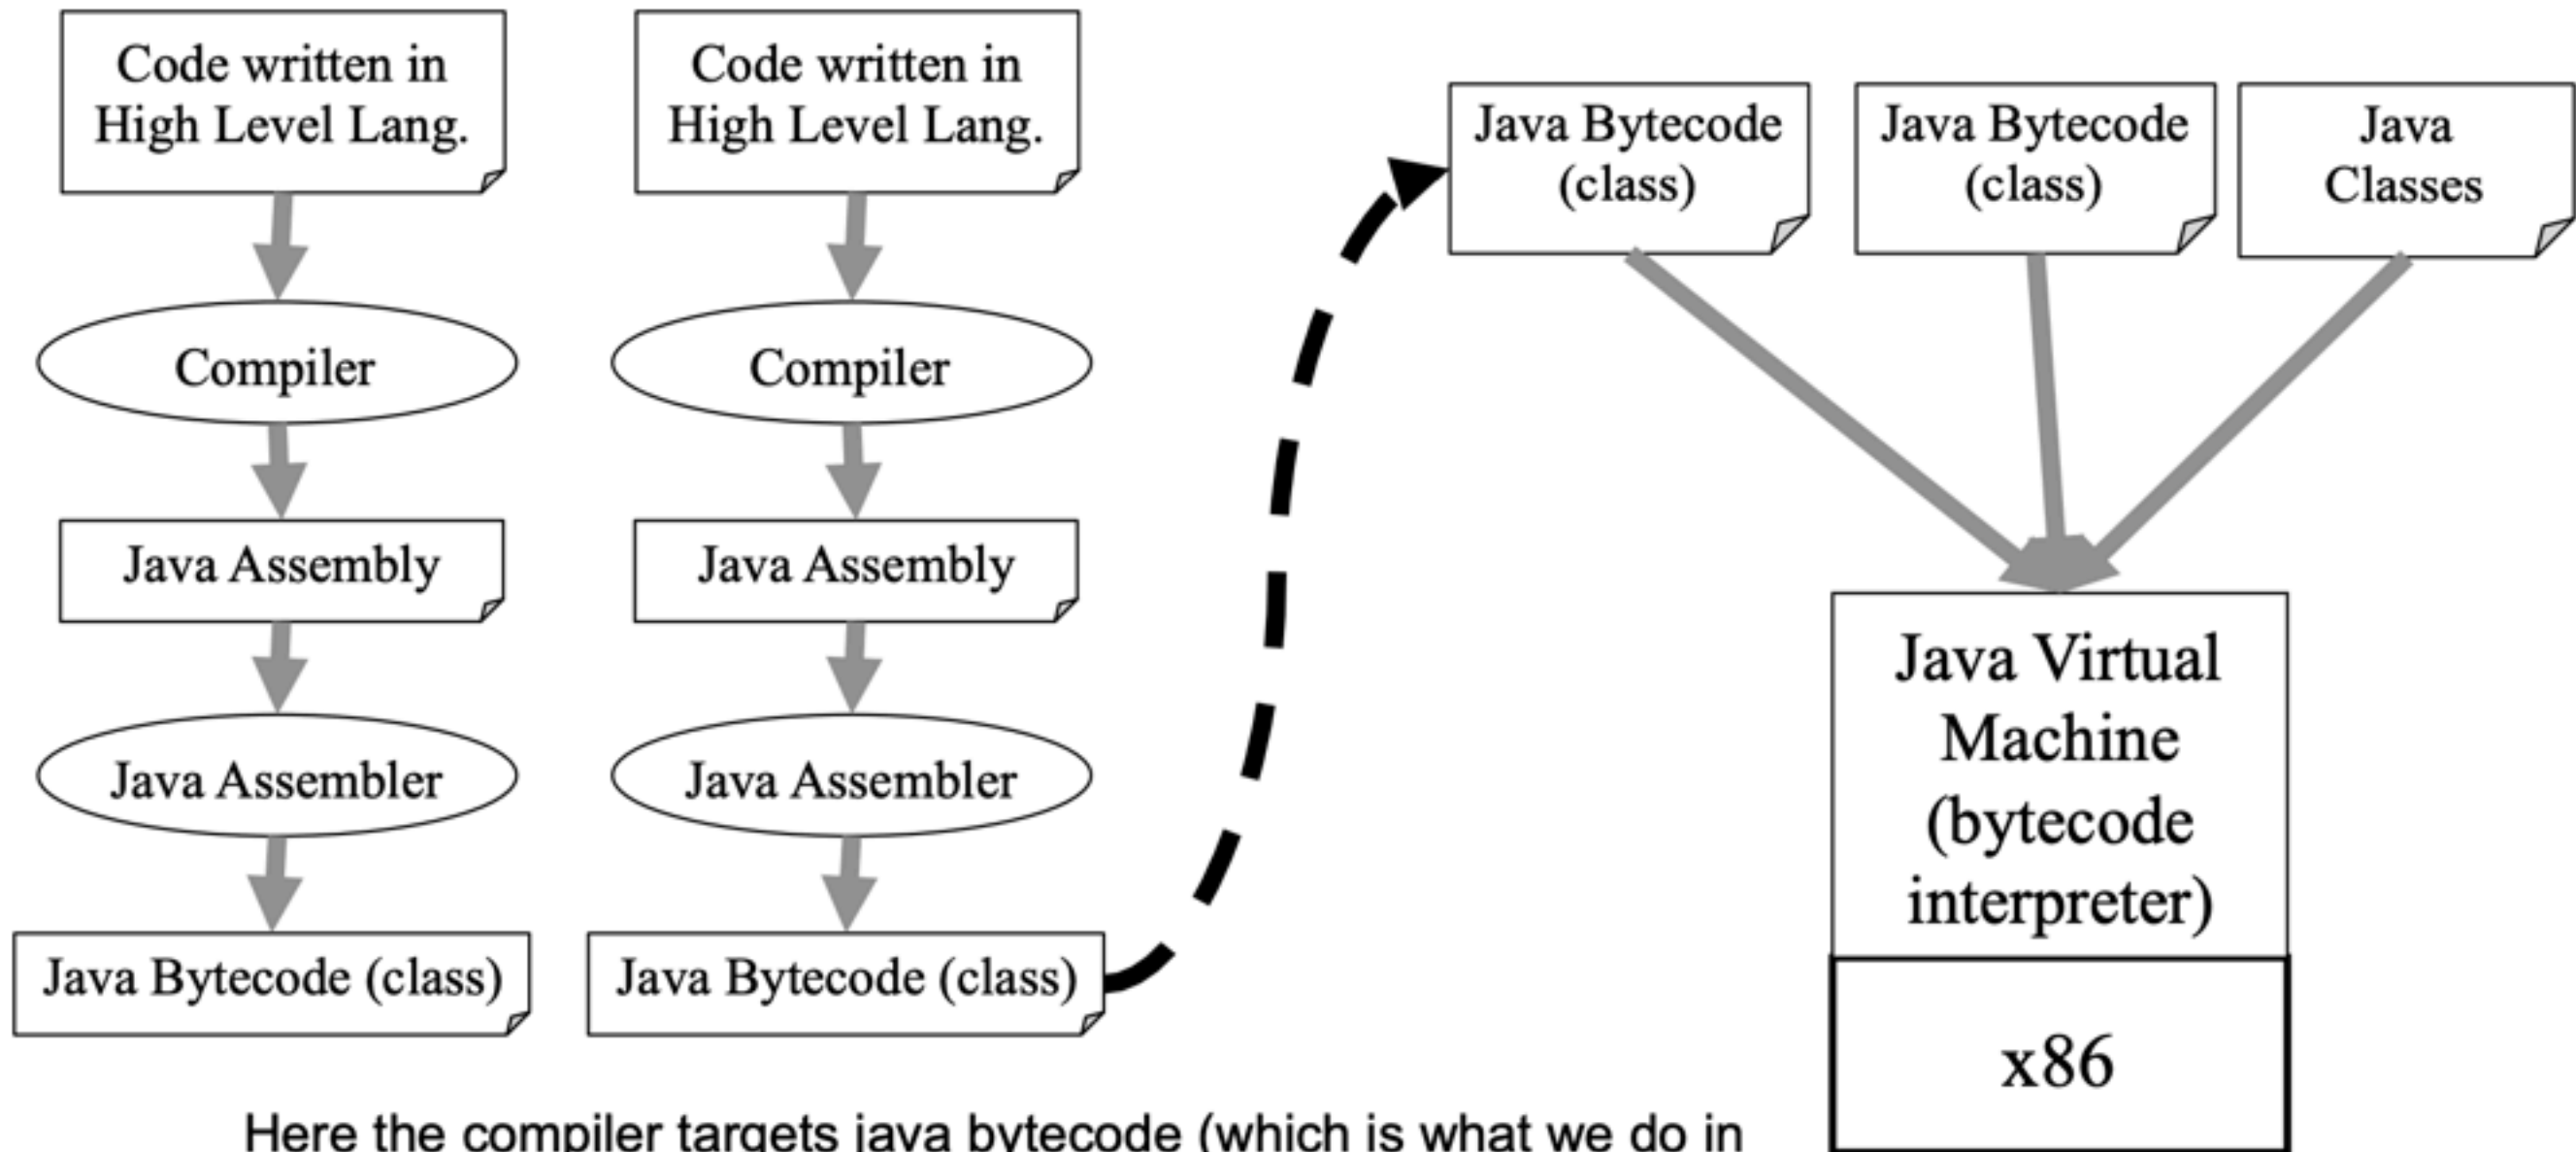
Code Generation

- To generate actual code that can run on a processor (such as gcc) or on a virtual machine (such as javac) we need to understand what code for each of these machines looks like.
- Rather than worry about the exact syntax of a given assembly language, we instead use a type of pseudo-assembly that is close to the underlying machine.
- In this class, we need to worry about 2 different types of code
 - Stack based code: Similar to the Java Virtual Machine
 - Register-based code: Similar to most processors (x86, Sparc, ARM)

x86 C Compiler



Java Compiler



Here the compiler targets java bytecode (which is what we do in this class) and the bytecode is then run on top of the Java Virtual Machine (JVM). The JVM really just interprets (simulates) the bytecode like any scripting language. Because of this, any java program compiled to bytecode is portable to any machine that someone has already ported the JVM too. No need to recompile.

Register-based Machine

- Each instruction can have at most three operands
- We have to break large statements into little operations that use temporary variables
 - $X=(2+3)+4$ turns into $T1=2+3; X=T1+4;$
- Temporary variables store the results at the internal nodes in the AST
- Assignments
 - $x := y$
 - $x := y \text{ op } z$ *op: binary arithmetic or logical operators*
 - $x := \text{op } y$ *op: unary operators (minus, negation, integer to float)*
- Branch
 - $\text{goto } L$ *execute the statement with labeled L next*
- Conditional Branch
 - $\text{if } x \text{ relop } y \text{ goto } L$ *relop: <, =, <=, >=, ==, !=*
 - if the condition holds, we execute statement labeled L next
 - if the condition does not hold, we execute the statement following this statement next

Register-based Machine

```
if (x < y)
    x = 5*y + 5*y/3;
else
    y = 5;
x = x + y;
```

Temporaries: temporaries correspond to the internal nodes of the syntax tree

Variables can be represented with their locations in the symbol table

```
if x < y goto L1
goto L2
L1: t1 := 5 * y
    t2 := 5 * y
    t3 := t2 / 3
    x := t1 + t2
    goto L3
L2: y := 5
L3: x := x + y
```

- Three-address code instructions can be represented as an array of
quadruples: operation, argument1, argument2, result
triples: operation, argument1, argument2
(each triple implicitly corresponds to a temporary)

Stack-based Machine

- Stack based code uses the stack to store temporary variables
- When we evaluate an expression $(E+E)$, it will take its arguments off the stack, add them together and put the result back on the stack.
- $(2+3)+4$ will *push 2; push 3; add; push 4; add*
- The machine code for this is a bit more ugly but the code is actually easier to generate because we do not need to handle temporary variables.

Stack-based Machine

```
if (x < y)
    x = 5*y + 5*y/3;
else
    y = 5;
x = x+y;
```

pushes the value
at the location x to
the stack

```
load x
load y
iflt L1
goto L2
L1: push 5
load y
multiply
push 5
load y
multiply
push 3
divide
add
store x
goto L3
L2: push 5
store y
L3: load x
load y
add
store x
```

pops the top
two elements and
compares them

pops the top two
elements, multiplies
them, and pushes the
result back to the stack

stores the value at the
top of the stack to the
location x

JVM: A stack machine

- JVM interpreter executes the bytecode on different machines
- JVM has an operand stack which we use to evaluate expressions
- JVM provides 65,535 local variables for each method
The local variables are like registers so we do not have to worry about register allocation
- Each local variable in JVM is denoted by a number between 0 and 65535 (x and y in the example will be assigned unique numbers)

Stack-based v.s. Register-based

- Register-Based code:
 - Good - Compact representation
 - Good - “Self contained”: has inputs, outputs, and operation all in one “instruction”
 - Bad - Requires lots of temporary variables
 - Bad - Temporary variables have to be handled explicitly
- Stack Based Code:
 - Good – No temporaries, everything is kept on the stack
 - Good – It is easy to generate code for this
 - Bad – Requires more instructions to do the same thing


Expressions

- We are targeting a register-based machine
- We need to evaluate expressions assuming a very limited set of available registers (No register allocation).
- To generate code for an expression we will do a recursive traversal in post-order (that is, visit the children first, then generate code for the parent).

Expressions

- Let's start with a simple example expression: $(1 + 2) * (3 - 4)$

```
call generate_aexp(* node, left):  
  call generate_aexp(+ node, left):  
    call generate_aexp(1 node, left):  
      emit "mov 1 LEFT_REG"  
    call generate_aexp(2 node, right):  
      emit "mov 2 RIGHT_REG"  
    emit "add RIGHT_REG LEFT_REG"  
  call generate_aexp(- node, right):  
    call generate_aexp(3 node, left):  
      emit "mov 3 LEFT_REG"  
    call generate_aexp(4 node, right):  
      emit "mov 4 RIGHT_REG"  
    emit "sub RIGHT_REG LEFT_REG" "mov LEFT_REG RIGHT_REG"  
  emit "mul RIGHT_REG LEFT_REG"
```




```
mov 1 LEFT_REG  
mov 2 RIGHT_REG  
add RIGHT_REG LEFT_REG  
mov 3 LEFT_REG  
mov 4 RIGHT_REG  
sub RIGHT_REG LEFT_REG  
mov LEFT_REG RIGHT_REG  
mul RIGHT_REG LEFT_REG
```

What is the problem?

Expressions

- We have to create memory locations to hold temporary values during expression evaluation.

```
call generate_aexp(* node, tmp_num = 0):  
  call generate_aexp(+ node, tmp_num = 1):  
    call generate_aexp(1 node, tmp_num = 2):  
      emit "mov 1 RESULT_REG"  
      insert _tmp1 into symbol table  
      emit "store RESULT_REG [_tmp1]"  
      call generate_aexp(2 node, tmp_num = 2):  
        emit "mov 2 RESULT_REG"  
        emit "ld [_tmp1] OTHER_REG" "add OTHER_REG RESULT_REG"  
        remove _tmp1 from symbol table  
      insert _tmp0 into symbol table  
      emit "store RESULT_REG [_tmp0]"  
    call generate_aexp(- node, tmp_num = 1):  
      insert _tmp1 into symbol table  
      call generate_aexp(3 node, tmp_num = 2):  
        emit "mov 3 RESULT_REG"  
        emit "store RESULT_REG [_tmp1]"  
      call generate_aexp(4 node, tmp_num = 2):  
        emit "mov 4 RESULT_REG"  
        emit "ld [_tmp1] OTHER_REG" "sub RESULT_REG OTHER_REG" "mov OTHER_REG RESULT_REG"  
        remove _tmp1 from symbol table  
      emit "ld [_tmp0] OTHER_REG" "mul OTHER_REG RESULT_REG"  
      remove _tmp0 from symbol table
```



```
mov 1 RESULT_REG  
store RESULT_REG [_tmp1]  
mov 2 RESULT_REG  
ld [_tmp1] OTHER_REG  
add OTHER_REG RESULT_REG  
store RESULT_REG [_tmp0]  
mov 3 RESULT_REG  
store RESULT_REG [_tmp1]  
mov 4 RESULT_REG  
ld [_tmp1] OTHER_REG  
sub RESULT_REG OTHER_REG  
mov OTHER_REG RESULT_REG  
ld [_tmp0] OTHER_REG  
mul OTHER_REG RESULT_REG
```

Let's Generalize It

- Let's generalize the algorithm for arbitrary arithmetic expressions

```
generate_aexp(AST* node, int tmp_num = 0) {
    if (node is a constant number <n>) { emit "mov <n> RESULT_REG";
return; }
    if (node is a variable <x>) { emit "ld [x] RESULT_REG"; return; }

    // node must be one of +,-,*
    generate_aexp(node->left, tmp_num+1);
    insert _tmp<tmp_num> into symbol table;
    emit "store RESULT_REG [_tmp<tmp_num>]";
    generate_aexp(node->right, tmp_num+1);
    emit "ld [_tmp<tmp_num>] OTHER_REG";

    // left-hand value is in OTHER_REG, right-hand value is in RESULT_REG
    if (node is +) { emit "add OTHER_REG RESULT_REG"; return; }
    if (node is -) { emit "sub RESULT_REG OTHER_REG"; emit "mov OTHER_REG
RESULT_REG"; return; }
    emit "mul OTHER_REG RESULT_REG";

    remove _tmp<tmp_num> from symbol table;
}
```

Assignment

- We evaluate the right-hand side using `generate_aexp`, which puts the result in `RESULT_REG`, then store the result to the memory location for the left-hand side variable.


```
generate_assign(lhs, rhs) {  
    generate_aexp(rhs);  
    emit "store RR [lhs]";  
}
```

Non-nested Conditionals

- Conditionals without a nested scope

```
if (x < 2) { x := 1; } else { x := 2; }
```

```
generate_if(node) {  
  <n> = fresh index;  
  generate_rexp(node->guard);  
  emit "cmp 0 RESULT_REG";  
  emit "jmpe IF_FALSE_<n>";  
  generate_block(node->true_branch);  
  emit "jmp IF_END_<n>";  
  emit "IF_FALSE_<n>:";  
  generate_block(node->false_branch);  
  emit "IF_END_<n>:";  
}
```



```
ld [x] RR  
store RR [_tmp0]  
mov 2 RR  
ld [_tmp0] OR  
cmp RR OR  
setlt RR  
cmp 0 RR  
jmpe IF_FALSE_0  
mov 1 RR  
store RR [x]  
jmp IF_END_0  
IF_FALSE_0:  
  mov 2 RR  
  store RR [x]  
IF_END_0:
```


Nested Conditionals

- When the code generator enters that new scope:
 - see how many declared variables there are
 - adjust the stack pointer accordingly
 - initialize the new memory locations to 0
 - update symbol table to map the newly declared variables to their offsets
- and when we leave the new scope we need to adjust things back the way they were:
 - reset the stack pointer to its old position
 - restore the symbol table to its old value

Nested Conditionals

```
generate_block(node) {  
    old_symbol_table = symbol_table;  
    stack_size = node->num_declared_variables * 4; // because 4-byte  
    integers  
    emit "sub <stack_size> STACK_REG";  
    insert_in_symbol_table(symbol_table, node->declared_variables);  
    for each var in node->declared_variables { emit "store 0 [var]"; }  
    .  
    .  
    .  
    emit "add <stack_size> STACK_REG";  
    symbol_table = old_symbol_table;  
}
```

While Loops

- Conditionals without a nested scope

```
while (x < 3) { x := x + 1; }
```



```
generate_while(node) {  
  <n> = fresh index;  
  emit "WHILE_START_<n>:";  
  generate_rexp(node->guard);  
  emit "cmp 0 RESULT_REG";  
  emit "jmpe WHILE_END_<n>";  
  generate_block(node->body);  
  emit "jmp WHILE_START_<n>";  
  emit "WHILE_END_<n>:";  
}
```

```
WHILE_START_0:  
  ld [x] RR  
  store RR [_tmp0]  
  mov 3 RR  
  ld [x] OR  
  cmp RR OR  
  setlt RR  
  cmp 0 RR  
  jmpe WHILE_END_0  
  ld [x] RR  
  store RR [_tmp0]  
  mov 1 RR  
  ld [_tmp0] OR  
  add OR RR  
  store RR [x]  
  jmp WHILE_START_0  
WHILE_END_0:
```

OCaml vs LLVM

“let - in” and
OCaml-style identifiers:

```
let tmp1 = add 3L 4L in
```

- OCaml-style “let-rec”
and functions for blocks:

```
let rec entry () =  
  let tmp1 = ...  
and foo () =  
  let tmp2 = ...
```

- OCaml-style global variables:

```
let varX = ref 0L
```

Omits let/in and prefixes local
identifiers with %:

```
%tmp1 = add i64 3, i64 4
```

- Uses lighter-weight colon notation:

```
entry:
```

```
  %tmp1 = ...
```

```
foo:
```

```
  %tmp2 = ...
```

- Prefixes globals with @

```
define @X = i64 0
```

Example LLVM Code

```
#include <stdio.h>
#include <stdint.h>

int64_t factorial(int64_t n) {
    int64_t acc = 1;
    while (n > 0) {
        acc = acc * n;
        n = n - 1;
    }
    return acc;
}
```



```
define @factorial(%n) {
    %1 = alloca
    %acc = alloca
    store %n, %1
    store 1, %acc
    br label %start

start:
    %3 = load %1
    %4 = icmp sgt %3, 0
    br %4, label %then, label %else

then:
    %6 = load %acc
    %7 = load %1
    %8 = mul %6, %7
    store %8, %acc
    %9 = load %1
    %10 = sub %9, 1
    store %10, %1
    br label %start

else:
    %12 = load %acc
    ret %12
}
```

Real LLVM

- Decorates values with type information

i64

i64*

i1

- Permits numeric identifiers
- Has alignment annotations
- Keeps track of entry edges for each block:
preds = %5, %0

```
; Function Attrs: nounwind ssp
define i64 @factorial(i64 %n) #0 {
    %1 = alloca i64, align 8
    %acc = alloca i64, align 8
    store i64 %n, i64* %1, align 8
    store i64 1, i64* %acc, align 8
    br label %2

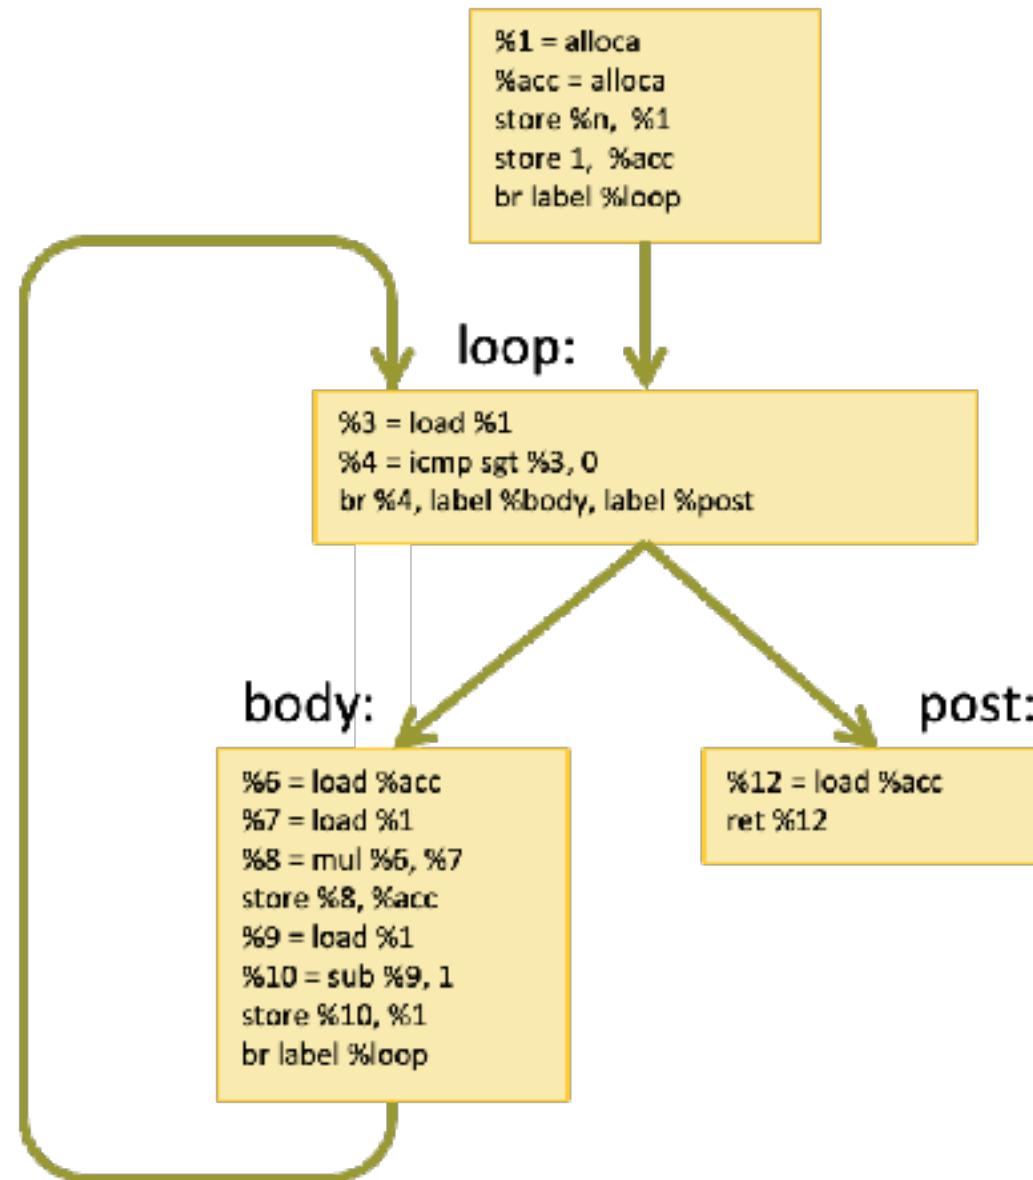
; <label>:2                ; preds = %5, %0
    %3 = load i64* %1, align 8
    %4 = icmp sgt i64 %3, 0
    br i1 %4, label %5, label %11

; <label>:5                ; preds = %2
    %6 = load i64* %acc, align 8
    %7 = load i64* %1, align 8
    %8 = mul nsw i64 %6, %7
    store i64 %8, i64* %acc, align 8
    %9 = load i64* %1, align 8
    %10 = sub nsw i64 %9, 1
    store i64 %10, i64* %1, align 8
    br label %2

; <label>:11               ; preds = %2
    %12 = load i64* %acc, align 8
    ret i64 %12
}
```

Example CFG

```
define @factorial(%n) {
```



```
}
```

LL Basic Blocks and Control-Flow Graphs

- LLVM enforces (some of) the basic block invariants syntactically.
- Representation in OCaml:

```
type block = {  
    insns : (uid * insn) list;  
    term : (uid * terminator)  
}
```

- A *control flow graph* is represented as a list of labeled basic blocks with these invariants:
 - No two blocks have the same label
 - All terminators mention only labels that are defined among the set of basic blocks
 - There is a distinguished, unlabeled, entry block:

```
type cfg = block * (lbl * block) list
```


LL Storage Model: Locals

- Several kinds of storage:
 - Local variables (or temporaries): `%uid`
 - Global declarations (e.g., for string constants): `@gid`
 - Abstract locations: references to (stack-allocated) storage created by the `alloca` instruction
 - Heap-allocated structures created by external calls (e.g., to `malloc`)
- Local variables:
 - Defined by the instructions of the form `%uid = ...`
 - Must satisfy the *static single assignment* invariant
 - *Each `%uid` appears on the left-hand side of an assignment only once in the entire control flow graph.*
 - The value of a `%uid` remains unchanged throughout its lifetime
 - Analogous to “`let %uid = e in ...`” in OCaml
- Intended to be an abstract version of machine registers.

LL Storage Model: alloca

- `alloca` instruction allocates stack space and returns a reference to it.
 - The returned reference is stored in local:
`%ptr = alloca type`
 - The amount of space allocated is determined by the type
- The contents of the slot are accessed via the load and store instructions:

```
%acc = alloca i64           ; allocate a storage slot
store i64 341, i64* %acc    ; store the integer value 341
%x = load i64, i64* %acc    ; load the value 341 into %x
```

- Gives an abstract version of stack slots

References

- llvm.org
- llvmLite: <https://www.cs.princeton.edu/courses/archive/spring19/cos320/hw/llvmlite.shtml>